

1 **Redistributing Interdependence in Organizational Knowledge Work: Lessons**
2 **from Deploying an LLM Mediator in Research Labs**
3

4 **ANONYMOUS AUTHOR(S)**
5

6 Large Language Models (LLMs) are increasingly embedded in collaborative workflows, yet their structural effects on the relationships
7 between human stakeholders in data-intensive knowledge work remain underexplored. In this position paper, we reinterpret findings
8 from a month-long field deployment of an LLM-based chatbot that mediates organizational memory across four university research
9 labs (N=21) through the lens of Interdependence Theory. Our analysis reveals two key dynamics. First, the LLM redistributes the
10 *dependence structure* between students and lab directors: while students gain autonomous access to institutional knowledge, directors
11 lose visibility into knowledge gaps, shifting bilateral dependence toward unilateral dependence. This redistribution is moderated by
12 organizational culture, amplifying *mutual responsiveness* in psychologically safe environments while dampening it where students
13 default to private interaction. Second, the system fails to support the *transformation of motivation* needed for collaborative knowledge
14 stewardship: students consistently avoid contributing to documentation due to role perceptions and temporal asymmetry between
15 individual costs and collective benefits. We derive design implications including privacy-preserving awareness mechanisms, graduated
16 contribution pathways, and designing for the LLM’s dual role as a boundary object across stakeholder groups.
17
18

19
20 CCS Concepts: • **Human-centered computing → Collaborative and social computing systems and tools.**
21

22
23 Additional Key Words and Phrases: human–LLM collaboration, interdependence theory, organizational memory, chatbot, knowledge
24 management
25
26
27

28 **1 Introduction**
29

30 Modern data and knowledge workflows are rarely solo efforts. They bring together people with different expertise,
31 vocabularies, and degrees of trust in automation [1, 9]. University research labs exemplify this dynamic: lab directors
32 maintain institutional knowledge, while students must locate, interpret, and contribute to that knowledge as part of
33 their academic work [14]. These workflows are inherently collaborative and data-intensive, involving the ongoing
34 collection, curation, and retrieval of organizational knowledge distributed across documents, conversations, and human
35 memory [2, 20]. At the same time, Large Language Models (LLMs) are increasingly embedded in tools for data processing,
36 knowledge retrieval, and documentation [10, 13]. While prior work has shown that AI intermediaries can reshape group
37 dynamics [3, 8, 17], how these reconfigurations alter the structural dependencies among stakeholders in collaborative
38 knowledge work remains less understood.
39
40

41 In this position paper, we draw on findings from a month-long field deployment of an LLM-based chatbot designed
42 to mediate organizational memory in university research labs [12]. Rather than treating the deployment solely as a
43 systems evaluation, we revisit its findings through the lens of *Interdependence Theory* [7, 19], a framework from social
44 psychology that analyzes how actors’ outcomes are mutually contingent on one another’s actions. We argue that
45 Interdependence Theory offers a productive lens for understanding human–LLM collaboration in knowledge work
46 because it foregrounds questions that purely technical evaluations often overlook: *How does the introduction of an LLM*
47 *change who depends on whom? What happens to mutual responsiveness when an AI mediates communication? How do*
48 *shared outcomes shift when individual actors can bypass collective channels?* We present an interdependence-theoretic
49 analysis of these findings and derive design implications for LLM-mediated collaborative systems.
50
51

53 2 Background

54 2.1 The Deployed System and Key Findings

55 The system analyzed in this paper is an LLM-based chatbot embedded in Slack that supports organizational memory
 56 in university research labs [12]. It provides document-grounded Q&A via retrieval-augmented generation [13], peer
 57 knowledge sharing with optional anonymity, conversational knowledge extraction, and director-approved document
 58 updates. All documentation changes require explicit director approval, ensuring that the director retains final authority
 59 over the organizational memory.

60 The system was deployed for one month across four university research labs (N=21). Participants submitted 107
 61 questions, 42% of which were answered using existing documentation. Lab directors made 38 documentation commits,
 62 adding 1,994 words to the organizational memory. The key findings that motivate our analysis include: (a) a *privacy-*
 63 *awareness tension*, where the majority of student questions were asked through private direct messages, limiting
 64 directors' visibility into knowledge gaps; (b) *culture-dependent usage*, where one lab's culture of psychological safety led
 65 to predominantly public questioning and significantly more documentation updates; and (c) *persistent barriers to student*
 66 *contribution*, where students hesitated to suggest documentation updates despite the system's collaborative features.
 67

72 2.2 Interdependence Theory

73 Interdependence Theory, originating from the work of Kelley and Thibaut [19], provides a framework for analyzing
 74 situations in which two or more actors' outcomes are mutually contingent on one another's actions [7, 16]. The theory
 75 characterizes the *dependence structure* between actors, meaning the degree and direction of influence each party's actions
 76 have on the other's outcomes. It further examines *mutual responsiveness*, the extent to which actors attend and adapt to
 77 each other's needs, and *transformation of motivation*, the process by which actors shift from self-interested behavior
 78 ("given preferences") toward behavior that accounts for the partner's or group's welfare ("effective preferences"). The
 79 theory also considers *correspondence of outcomes*, whether actors' interests are aligned or in conflict. While originally
 80 formulated for human-human dyads, the theory's emphasis on structural features of interaction situations, rather than
 81 on cognitive capacities, makes it applicable to analyzing human-LLM configurations.
 82

86 3 An Interdependence-Theoretic Analysis

87 3.1 Redistribution of Dependence and Disrupted Mutual Responsiveness

88 Prior to the introduction of the LLM chatbot, the *dependence structure* in research labs was relatively straightforward.
 89 Students depended on lab directors for access to institutional knowledge (policies, procedures, implicit norms), and
 90 lab directors depended on students' questions to understand their knowledge needs and to maintain awareness of
 91 documentation gaps. This created a bilateral dependence structure: each party's outcomes (effective knowledge access
 92 for students; well-maintained documentation for directors) were contingent on the other's actions.
 93

94 The introduction of the LLM mediator fundamentally restructured this dependence. Students' dependence on directors
 95 for routine knowledge retrieval was partially displaced onto the system, as 42% of their questions could be answered
 96 directly from existing documentation. This displacement was welcomed by students, who described the chatbot as
 97 "a senior in the lab" or "a librarian," a dependable knowledge source that reduced the social cost of asking questions.
 98 However, this redistribution simultaneously attenuated directors' dependence on student questions as a source of
 99 awareness. When students obtained answers privately, the informational signal that directors relied upon (observing
 100 what students asked and where they struggled) was no longer available. This asymmetric redistribution created a
 101

105 new interdependence structure in which students' outcomes (knowledge access) improved while directors' outcomes
106 (documentation awareness) were inadvertently diminished. In Interdependence Theory terms, this represents a shift
107 from bilateral to more unilateral dependence: the LLM, in absorbing one pathway of dependence, inadvertently severed
108 a feedback loop that sustained the collaborative maintenance of organizational knowledge.

109 Critically, this redistribution manifested differently depending on organizational culture. In one lab (Lab A), where
110 a culture of psychological safety led members to ask questions publicly, the LLM amplified *mutual responsiveness*:
111 directors could observe student needs, discussions fed into documentation updates, and Lab A's director made the
112 most commits (19). In the remaining three labs, where students predominantly used private direct messages, the LLM
113 dampened mutual responsiveness; one director expressed surprise upon learning how actively students had been
114 using the system privately. The same LLM system can either amplify or dampen mutual responsiveness depending on
115 the organizational culture in which it is embedded. From an interdependence-theoretic perspective, organizational
116 culture functions as a moderator of the situation structure [7], producing different effective situations depending on the
117 interpersonal norms, trust levels, and power dynamics of the group.
118
119
120
121

122 3.2 Failures of Transformation in Knowledge Contribution

123

124 The findings reveal systematic failures of what Interdependence Theory calls *transformation of motivation*, the process
125 by which actors shift from self-interested behavior ("given preferences") to prosocial behavior that considers the
126 group's welfare ("effective preferences"). Students recognized that contributing their knowledge to organizational
127 documentation would benefit future lab members, yet they consistently prioritized their given preferences: avoiding
128 the risk of documenting inaccurate information, protecting themselves from social evaluation, and minimizing the
129 effort of generalizing personal experiences. As one student noted, the concern was that contributed knowledge "would
130 stay in the database forever" and could "mislead all the students."

131 Notably, this reluctance persisted even though the system's design included a safeguard: all documentation updates
132 required explicit approval from the lab director before being applied. Students were aware of this review process, yet
133 the knowledge that a director would see and evaluate their contribution did not alleviate their hesitation; if anything, it
134 added a layer of social evaluation. More fundamentally, students perceived documentation maintenance as the director's
135 prerogative and responsibility rather than a shared endeavor. This framing positioned students as passive consumers of
136 organizational knowledge and directors as sole curators, reinforcing an asymmetric dependence structure in which
137 students depended on directors for both knowledge access and knowledge maintenance. Even when the system provided
138 technical pathways for student contribution, the perceived role boundary limited the transformation of motivation
139 necessary for collaborative knowledge stewardship.

140 This failure of transformation also reflected a structural property of the interdependence situation. Contributing to
141 documentation creates a temporal asymmetry: the costs (effort, risk of error, social exposure) are borne immediately by
142 the contributor, while the benefits accrue to future, often unknown, lab members. In Interdependence Theory terms, this
143 is a *low-correspondence* situation where the contributor's immediate interests diverge from the collective's long-term
144 interests. The chatbot's Knowledge Extraction feature partially addressed this by lowering the cost of contribution,
145 transforming ongoing conversations into documentation without requiring separate writing effort. However, the
146 psychological barriers persisted, suggesting that LLM-mediated systems must go beyond reducing technical friction
147 and actively address the role perceptions and social dynamics that inhibit prosocial transformation.

157 4 Design Implications

158 Our interdependence-theoretic analysis yields several design implications for LLM-mediated collaborative systems,
 159 particularly those operating in knowledge-intensive, data-rich organizational contexts.
 160

161 *Privacy-preserving interdependence mechanisms.* The privacy-awareness tension represents a fundamental design
 162 challenge [5]: how to maintain the informational interdependence that directors need for effective documentation
 163 maintenance without compromising students' autonomy and psychological safety. We recommend *aggregated awareness*
 164 *mechanisms*, such as periodic, anonymized summaries of question patterns, knowledge gaps, and documentation
 165 usage, that preserve the interdependence signal without exposing individual behavior. This approach maintains the
 166 structural interdependence between stakeholders while respecting individual preferences for privacy, aligning with
 167 what participants in the study explicitly suggested.
 168

169 *Graduated transformation support.* To address the observed failures of transformation [14], LLM-mediated systems
 170 should provide graduated pathways for knowledge contribution [11]. Rather than requiring contributors to immediately
 171 produce universally applicable documentation, systems should support a spectrum, from low-commitment annotations
 172 ("this worked for me") to collaborative refinement (peer discussion and director review) to formal documentation.
 173 Automatically extracting experiential knowledge from ongoing conversations and drafting documentation from them
 174 can lower the barrier from "authoring documentation" to "sharing what happened." Crucially, the design must also
 175 address role perceptions: when students view documentation as exclusively the director's domain, even well-designed
 176 contribution mechanisms may go unused. Systems could reframe contribution as a normal part of membership in the
 177 organization rather than an encroachment on managerial authority.
 178

179 *Designing for the LLM's dual role as boundary object.* The deployment revealed that users simultaneously experienced
 180 the LLM as a social agent—students described it as "a senior in the lab"—and as a communication medium through
 181 which directors broadcast documentation updates. This duality suggests that LLM mediators in organizational contexts
 182 function as boundary objects [18]: artifacts that maintain shared identity across different social worlds while serving
 183 each group's situated needs [4, 6, 15]. Designers should account for this dual role, ensuring that the same system
 184 can support private, low-stakes knowledge retrieval for individual users while also enabling collective knowledge
 185 maintenance and organizational awareness for managers.
 186

187 References

- 195 [1] Mark S. Ackerman, Juri Dachtera, Volkmar Pipek, and Volker Wulf. 2013. Sharing Knowledge and Expertise: The CSCW View of Knowledge
 196 Management. *Computer Supported Cooperative Work (CSCW)* 22, 4 (2013), 531–573. doi:[10.1007/s10606-013-9192-8](https://doi.org/10.1007/s10606-013-9192-8)
- 197 [2] Mark S. Ackerman and Christine Halverson. 2004. Organizational Memory as Objects, Processes, and Trajectories: An Examination of Organizational
 198 Memory in Use. *Computer Supported Cooperative Work (CSCW)* 13, 2 (2004), 155–189. doi:[10.1023/B:COSU.0000045805.77534.2a](https://doi.org/10.1023/B:COSU.0000045805.77534.2a)
- 199 [3] Chun-Wei Chiang, Zhuoran Lu, Zhuoyan Li, and Ming Yin. 2024. Enhancing AI-Assisted Group Decision Making through LLM-Powered Devil's
 200 Advocate. In *Proceedings of the 29th International Conference on Intelligent User Interfaces (IUI '24)*. Association for Computing Machinery, New York,
 201 NY, USA, 103–119. doi:[10.1145/3640543.3645199](https://doi.org/10.1145/3640543.3645199)
- 202 [4] John Joon Young Chung and Eytan Adar. 2023. Artinter: AI-powered Boundary Objects for Commissioning Visual Arts. In *Proceedings of the 2023
 203 ACM Designing Interactive Systems Conference (DIS '23)*. Association for Computing Machinery, 1997–2018. doi:[10.1145/3563657.3595961](https://doi.org/10.1145/3563657.3595961)
- 204 [5] Paul Dourish and Victoria Bellotti. 1992. Awareness and Coordination in Shared Workspaces. In *Proceedings of the 1992 ACM Conference on
 205 Computer-supported Cooperative Work - CSCW '92*. ACM Press, Toronto, Ontario, Canada, 107–114. doi:[10.1145/143457.143468](https://doi.org/10.1145/143457.143468)
- 206 [6] Jose A. Guridi, Angel Hsing-Chi Hwang, Duarte Santo, Maria Goula, Cristobal Cheyre, Lee Humphreys, and Marco Rangel. 2025. From Fake
 207 Perfects to Conversational Imperfects: Exploring Image-Generative AI as a Boundary Object for Participatory Design of Public Spaces. *Proc. ACM
 208 Hum.-Comput. Interact.* 9, 2 (2025). doi:[10.1145/3710912](https://doi.org/10.1145/3710912)

- 209 [7] Harold H Kelley, John G Holmes, Norbert L Kerr, Harry T Reis, Caryl E Rusbult, and Paul AM Van Lange. 2003. An atlas of interpersonal situations.
210 (2003).
- 211 [8] Soomin Kim, Jinsu Eun, Joseph Seering, and Joonhwan Lee. 2021. Moderator Chatbot for Deliberative Discussion: Effects of Discussion Structure
212 and Discussant Facilitation. *Proc. ACM Hum.-Comput. Interact.* 5, CSCW1 (April 2021), 87:1–87:26. doi:[10.1145/3449161](https://doi.org/10.1145/3449161)
- 213 [9] Kalle Koivisto and Toni Taipalus. 2025. Pitfalls in Effective Knowledge Management: Insights from an International Information Technology
214 Organization. *International Journal of Knowledge Management Studies* 16, 1 (2025), 1–26. arXiv:[2304.07737](https://arxiv.org/abs/2304.07737) [cs] doi:[10.1504/IJKMS.2025.146083](https://doi.org/10.1504/IJKMS.2025.146083)
- 215 [10] Philippe Laban, Jesse Vig, Marti Hearst, Caiming Xiong, and Chien-Sheng Wu. 2024. Beyond the Chat: Executable and Verifiable Text-Editing
216 with LLMs. In *Proceedings of the 37th Annual ACM Symposium on User Interface Software and Technology (UIST '24)*. Association for Computing
217 Machinery, New York, NY, USA, 1–23. doi:[10.1145/3654777.3676419](https://doi.org/10.1145/3654777.3676419)
- 218 [11] Jean Lave and Etienne Wenger. 1991. *Situated Learning: Legitimate Peripheral Participation*. Cambridge University Press.
- 219 [12] Sangwook Lee, Adnan Abbas, Yan Chen, Young-Ho Kim, and Sang Won Lee. 2025. CHOIR: A Chatbot-mediated Organizational Memory Leveraging
220 Communication in University Research Labs. *arXiv preprint arXiv:2509.20512* (2025).
- 221 [13] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim
222 Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2021. Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. arXiv:[2005.11401](https://arxiv.org/abs/2005.11401) [cs]
223 doi:[10.48550/arXiv.2005.11401](https://doi.org/10.48550/arXiv.2005.11401)
- 224 [14] Lynne Markus. 2001. Toward a Theory of Knowledge Reuse: Types of Knowledge Reuse Situations and Factors in Reuse Success. *Journal of
Management Information Systems* 18, 1 (May 2001), 57–93. doi:[10.1080/07421222.2001.11045671](https://doi.org/10.1080/07421222.2001.11045671)
- 225 [15] Julien Porquet, Sitong Wang, and Lydia B. Chilton. 2025. Copying style, Extracting value: Illustrators' Perception of AI Style Transfer and its
226 Impact on Creative Labor. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems (CHI '25)*. Association for Computing
227 Machinery. doi:[10.1145/3706598.3713854](https://doi.org/10.1145/3706598.3713854)
- 228 [16] Caryl E Rusbult and Paul AM Van Lange. 2003. Interdependence, interaction, and relationships. *Annual review of psychology* 54, 1 (2003), 351–375.
- 229 [17] Donghoon Shin, Soomin Kim, Ruoxi Shang, Joonhwan Lee, and Gary Hsieh. 2023. IntroBot: Exploring the Use of Chatbot-assisted Familiarization
230 in Online Collaborative Groups. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems (CHI '23)*. Association for Computing
231 Machinery, New York, NY, USA, 1–13. doi:[10.1145/3544548.3580930](https://doi.org/10.1145/3544548.3580930)
- 232 [18] Susan Leigh Star and James R. Griesemer. 1989. Institutional Ecology, 'Translations' and Boundary Objects: Amateurs and Professionals in Berkeley's
233 Museum of Vertebrate Zoology, 1907–39. *Social Studies of Science* 19, 3 (1989), 387–420. doi:[10.1177/030631289019003001](https://doi.org/10.1177/030631289019003001)
- 234 [19] John W Thibaut. 2017. *The social psychology of groups*. Routledge.
- 235 [20] James P. Walsh and Gerardo Rivera Ungson. 1991. Organizational Memory. *The Academy of Management Review* 16, 1 (1991), 57–91. doi:[10.2307/258607](https://doi.org/10.2307/258607)
- 236
- 237
- 238
- 239
- 240
- 241
- 242
- 243
- 244
- 245
- 246
- 247
- 248
- 249
- 250
- 251
- 252
- 253
- 254
- 255
- 256
- 257
- 258
- 259
- 260